

Telco Customer Churn Analysis

January 4, 2026

1 Telco Customer Churn Analysis — End-to-End Case Study

1.1 Project Objective

This project analyzes customer churn for a telecommunications company.

Churn Definition A customer is considered churned if:

- Churn = "Yes" → Customer has left the company
- Churn = "No" → Customer is still active

Churn is a critical metric because:

- Retaining customers costs less than acquiring new ones
- Churn reduces recurring revenue & lifetime value
- High churn signals dissatisfaction & competition pressure

1.2 Project Goals

This project aims to:

1. Understand churn behaviour and key drivers
2. Compare churn vs non-churn customer groups
3. Engineer behavioural & lifecycle features
4. Train churn prediction models
5. Evaluate model performance
6. Provide actionable business recommendations

1.3 Dataset Description

Each row represents a customer record.

Key attributes include:

- Demographics
- Account information
- Subscription plans
- Billing details

- Service usage
- Churn outcome

The dataset is ideal for an **end-to-end churn analytics case study**.

1.3.1 Data Cleaning & Preparation

Import Libraries & Load Data

```
[1]: import pandas as pd
import numpy as np

df = pd.read_csv("WA_Fn-UseC_-Telco-Customer-Churn 2 -_
↳WA_Fn-UseC_-Telco-Customer-Churn 2.csv")

df.head()
```

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService	OnlineSecurity	... DeviceProtection	TechSupport	StreamingTV	StreamingMovies	Contract	PaperlessBilling	PaymentMethod	MonthlyCharges	TotalCharges	Churn
0	7590-VHVEG	Female	0	Yes	No	1	No	No phone service	DSL	No	...	No	No	No	Month-to-month	Yes	Electronic check	29.85	29.85	No
1	5575-GNVDE	Male	0	No	No	34	Yes	No	DSL	Yes	...	Yes	No	No	One year	No	Mailed check	56.95	1889.50	No
2	3668-QPYBK	Male	0	No	No	2	Yes	No	DSL	Yes	...	No	No	Month-to-month	Yes	Mailed check	53.85	108.15	Yes	
3	7795-CFOCW	Male	0	No	No	45	No	No	DSL	Yes	...	Yes	No	One year	No	Bank transfer (automatic)	42.30	1840.75	No	
4	9237-HQITU	Female	0	No	No	2	Yes	No	Fiber optic	No	...	No	No	Month-to-month	Yes	Electronic check	70.70	151.65	Yes	

[5 rows x 21 columns]

Dataset Structure & Summary

Inspect column types, shape, and completeness.

[3]: df.shape, df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 21 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   customerID      7043 non-null   object  
 1   gender          7043 non-null   object  
 2   SeniorCitizen   7043 non-null   int64  
 3   Partner         7043 non-null   object  
 4   Dependents     7043 non-null   object  
 5   tenure          7043 non-null   int64  
 6   PhoneService    7043 non-null   object  
 7   MultipleLines   7043 non-null   object  
 8   InternetService 7043 non-null   object  
 9   OnlineSecurity  7043 non-null   object  
 10  OnlineBackup    7043 non-null   object  
 11  DeviceProtection 7043 non-null   object  
 12  TechSupport     7043 non-null   object  
 13  StreamingTV     7043 non-null   object  
 14  StreamingMovies 7043 non-null   object  
 15  Contract        7043 non-null   object  
 16  PaperlessBilling 7043 non-null   object  
 17  PaymentMethod    7043 non-null   object  
 18  MonthlyCharges  7043 non-null   float64 
 19  TotalCharges    7032 non-null   float64 
 20  Churn           7043 non-null   object  
dtypes: float64(2), int64(2), object(17)
memory usage: 1.1+ MB
```

[3]: ((7043, 21), None)

To check for missing values and ensure data consistency.

[5]: df.isna().sum()

```
[5]: customerID      0
gender          0
SeniorCitizen   0
Partner         0
Dependents     0
tenure          0
PhoneService    0
MultipleLines   0
InternetService 0
OnlineSecurity  0
```

```
OnlineBackup      0
DeviceProtection  0
TechSupport       0
StreamingTV       0
StreamingMovies   0
Contract          0
PaperlessBilling  0
PaymentMethod     0
MonthlyCharges    0
TotalCharges      11
Churn             0
dtype: int64
```

```
[7]: df['TotalCharges'] = df['TotalCharges'].fillna(0)
df['TotalCharges'] = pd.to_numeric(df['TotalCharges'], errors='coerce')
```

Numeric churn flag

```
[9]: df['Churn_flag'] = (df['Churn'] == 'Yes').astype(int)
```

```
[11]: %matplotlib inline
import matplotlib.pyplot as plt
import seaborn as sns

sns.set(style="whitegrid")
plt.rcParams["figure.figsize"] = (10,5)
```

1.3.2 Churn Distribution

To compute the churn ratio to understand the dataset balance.


```
[13]: churn_counts = df['Churn'].value_counts()

sns.barplot(x=churn_counts.index, y=churn_counts.values, palette="pastel")
plt.title("Customer Churn Distribution")
plt.ylabel("Number of Customers")
plt.xlabel("Churn Status")
plt.show()
```

C:\Users\mildr\AppData\Local\Temp\ipykernel_6848\1081618862.py:3: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(x=churn_counts.index, y=churn_counts.values, palette="pastel")
```



1.4 Exploratory Data Analysis (EDA)

Goal:

- Compare churn vs non-churn groups
- Identify high-risk churn segments
- Support feature engineering & modeling

1.4.1 Churn vs Contract Type

```
[15]: contract_churn = (
    df.groupby('Contract')['Churn_flag'].mean().reset_index()
)


sns.barplot(
    data=contract_churn,
    x='Contract',
    y='Churn_flag',
    palette="viridis"
)

plt.title("Churn Rate by Contract Type")
plt.ylabel("Churn Rate")
plt.xlabel("Contract Type")
plt.show()
```

C:\Users\mildr\AppData\Local\Temp\ipykernel_6848\1435280597.py:5: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(
```



Short-term contracts = **highest churn risk**

1.4.2 Churn vs Tenure Groups

```
[17]: bins = [0,12,24,48,60,72]
labels = ['0-1yr', '1-2yrs', '2-4yrs', '4-5yrs', '5-6yrs']

df['tenure_group'] = pd.cut(df['tenure'], bins=bins, labels=labels)

tenure_churn = (
    df.groupby('tenure_group')['Churn_flag'].mean().reset_index()
)

sns.barplot(
    data=tenure_churn,
    x='tenure_group',
    y='Churn_flag',
    palette="magma"
)


plt.title("Churn Rate by Tenure Group")
plt.ylabel("Churn Rate")
plt.xlabel("Customer Tenure Group")
plt.show()
```

```
C:\Users\mildr\AppData\Local\Temp\ipykernel_6848\1005993199.py:7: FutureWarning:  
The default of observed=False is deprecated and will be changed to True in a  
future version of pandas. Pass observed=False to retain current behavior or  
observed=True to adopt the future default and silence this warning.
```

```
df.groupby('tenure_group')['Churn_flag'].mean().reset_index()  
C:\Users\mildr\AppData\Local\Temp\ipykernel_6848\1005993199.py:10:  
FutureWarning:
```

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(
```



Retention is most critical during first 12 months

1.4.3 Churn vs Internet Service

```
[19]: internet_churn = (  
    df.groupby('InternetService')['Churn_flag'].mean().reset_index()  
)  
  
sns.barplot(  
    data=internet_churn,  
    x='InternetService',  
    y='Churn_flag',  
    palette="cubehelix"
```

```


)
plt.title("Churn Rate by Internet Service Type")
plt.ylabel("Churn Rate")
plt.xlabel("Internet Service")
plt.show()

```

C:\Users\mildr\AppData\Local\Temp\ipykernel_6848\1479068448.py:5: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(
```



Fiber users may be:

- . more price-sensitive
- . higher performance expectations

1.4.4 Churn vs Payment Method

```
[21]: payment_churn = (
    df.groupby('PaymentMethod')['Churn_flag'].mean().reset_index()
)

sns.barplot(
```

```

    data=payment_churn,
    x='PaymentMethod',
    y='Churn_flag',
    palette="coolwarm"
)


plt.title("Churn Rate by Payment Method")
plt.ylabel("Churn Rate")
plt.xlabel("Payment Method")
plt.xticks(rotation=45)
plt.show()

```

C:\Users\mildr\AppData\Local\Temp\ipykernel_6848\413570720.py:5: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

sns.barplot(



Auto-pay customers show higher retention stability

1.5 Feature Engineering

To create behavioural and lifecycle features.

Number of subscribed services

```
[23]: service_cols = [
    'PhoneService', 'MultipleLines', 'OnlineSecurity', 'OnlineBackup',
    'DeviceProtection', 'TechSupport', 'StreamingTV', 'StreamingMovies'
]

df['num_services'] = df[service_cols].apply(lambda x: (x == 'Yes').sum(), axis=1)
```

Tenure lifecycle bins

```
[25]: df['tenure_group'] = pd.cut(df['tenure'], bins=bins, labels=labels)
```

Final feature matrix

```
[27]: X = df.drop(columns=['customerID', 'Churn', 'Churn_flag'])
y = df['Churn_flag']
```

1.6 Train-Test Split & Preprocessing

Model Training Pipeline

Includes preprocessing + encoding + scaling.

```
[29]: from sklearn.model_selection import train_test_split
from sklearn.preprocessing import OneHotEncoder, StandardScaler
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline

cat_cols = X.select_dtypes(include=['object']).columns.tolist()
num_cols = X.select_dtypes(include=['int64', 'float64']).columns.tolist()

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)

preprocess = ColumnTransformer([
    ('cat', OneHotEncoder(drop='first', handle_unknown='ignore'), cat_cols),
    ('num', StandardScaler(), num_cols)
])
```

Logic Regression Model

```
[31]: from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report, confusion_matrix,
    roc_auc_score
```

```

log_model = Pipeline(steps=[
    ('prep', preprocess),
    ('clf', LogisticRegression(max_iter=1000))
])

log_model.fit(X_train, y_train)

y_pred_lr = log_model.predict(X_test)
y_prob_lr = log_model.predict_proba(X_test)[:,1]

print(confusion_matrix(y_test, y_pred_lr))
print(classification_report(y_test, y_pred_lr))
print("ROC-AUC:", roc_auc_score(y_test, y_prob_lr))

```

```

[[926 109]
 [165 209]]
      precision    recall   f1-score   support
          0       0.85     0.89     0.87     1035
          1       0.66     0.56     0.60      374

   accuracy                           0.81     1409
  macro avg       0.75     0.73     0.74     1409
weighted avg       0.80     0.81     0.80     1409

```

ROC-AUC: 0.8422666563331527

Random Forest Model

```
[33]: from sklearn.ensemble import RandomForestClassifier

rf_model = Pipeline(steps=[
    ('prep', preprocess),
    ('clf', RandomForestClassifier(
        n_estimators=300,
        random_state=42,
        min_samples_split=5,
        min_samples_leaf=2,
        class_weight='balanced'
    ))
])

rf_model.fit(X_train, y_train)

y_pred_rf = rf_model.predict(X_test)
y_prob_rf = rf_model.predict_proba(X_test)[:,1]
```

```

print(confusion_matrix(y_test, y_pred_rf))
print(classification_report(y_test, y_pred_rf))
print("ROC-AUC:", roc_auc_score(y_test, y_prob_rf))

```

```

[[856 179]
 [133 241]]

      precision    recall   f1-score   support

          0       0.87      0.83      0.85      1035
          1       0.57      0.64      0.61      374

   accuracy                           0.78      1409
  macro avg       0.72      0.74      0.73      1409
weighted avg       0.79      0.78      0.78      1409

```

ROC-AUC: 0.8366374744891368

Feature Importance (Random Forest)

```

[35]: ohe = rf_model.named_steps['prep'].named_transformers_['cat']
cat_features = ohe.get_feature_names_out(cat_cols)

feature_names = list(cat_features) + num_cols

importances = rf_model.named_steps['clf'].feature_importances_

fi = pd.Series(importances, index=feature_names).sort_values(ascending=False)
fi.head(15)

```

```

[35]: tenure                      0.175337
TotalCharges                   0.152870
MonthlyCharges                  0.119737
Contract_Two year               0.072586
InternetService_Fiber optic     0.060801
PaymentMethod_Electronic check   0.040468
num_services                     0.033757
Contract_One year                0.033624
OnlineSecurity_Yes                 0.028361
TechSupport_Yes                   0.022761
PaperlessBilling_Yes                0.021337
gender_Male                       0.019548
OnlineBackup_Yes                   0.017432
Partner_Yes                        0.016767
Dependents_Yes                     0.015592
dtype: float64

```

1.7 Post-Dictive Churn Analysis

We compare actual vs predicted churn at the segment level.

```
[37]: test_df = X_test.copy()
test_df['actual'] = y_test.values
test_df['pred'] = y_pred_rf
test_df['prob'] = y_prob_rf

test_df.head()
```

	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	\
437	Male	0	Yes	Yes	72	Yes	
2280	Female	1	No	No	8	Yes	
2235	Female	0	Yes	Yes	41	Yes	
4460	Male	0	Yes	No	18	Yes	
3761	Female	0	Yes	No	72	Yes	
	MultipleLines	InternetService	OnlineSecurity	OnlineBackup	...	\	
437	Yes	Fiber optic		Yes	Yes	...	
2280	Yes	Fiber optic		No	No	...	
2235	Yes		DSL	Yes	Yes	...	
4460	No	Fiber optic		No	No	...	
3761	Yes		DSL	Yes	Yes	...	
	Contract	PaperlessBilling		PaymentMethod	MonthlyCharges	\	
437	Two year	Yes	Credit card (automatic)		114.05		
2280	Month-to-month	Yes	Credit card (automatic)		100.15		
2235	One year	Yes	Credit card (automatic)		78.35		
4460	Month-to-month	No	Electronic check		78.20		
3761	Two year	Yes	Credit card (automatic)		82.65		
	TotalCharges	tenure_group	num_services	actual	pred	prob	
437	8468.20	5-6yrs		8	0	0.010978	
2280	908.55	0-1yr		5	0	0.850672	
2235	3211.20	2-4yrs		6	0	0.132907	
4460	1468.75	1-2yrs		3	0	0.402107	
3761	5919.35	5-6yrs		7	0	0.019593	

[5 rows x 24 columns]

Identify high-risk segment

```
[39]: segment = test_df[
    (test_df['Contract'] == 'Month-to-month') &
    (test_df['InternetService'] == 'Fiber optic') &
    (test_df['tenure'] <= 12)
]

segment[['actual', 'pred']].mean()
```

```
[39]: actual    0.661290
      pred     0.978495
      dtype: float64
```

2 Business Recommendations

2.0.1 High-Risk Churn Segments

- Month-to-month customers
- Tenure below 12 months
- Fiber-optic users
- Electronic-check payment users

2.1 Recommended Retention Actions

2.1.1 1. Convert monthly contracts → 1-year plans

- loyalty rewards
- contract upgrade discounts

2.1.2 2. First-year onboarding program

- welcome support calls
- quarterly check-ins
- service education touchpoints

2.1.3 3. Fiber-customer experience improvement

- performance investigation
- premium support tier
- bundled feature offers

2.1.4 4. Move customers to auto-payment

- incentive for bank transfer/card autopay

2.1.5 5. Promote value-added services

Customers with: - TechSupport - OnlineSecurity

show **lower churn probability**

2.2 Final Conclusion

The analysis shows that churn is driven by:

- contract flexibility
- early-lifecycle risk
- billing method
- internet plan type
- service value perception

The Random Forest model provides:

- strong recall for churners
- reliable churn risk segmentation

This supports **data-driven retention strategy design**.

[]: